

Generative AI Training

City of Garden Grove Employees

1. Why This Training

Reality: AI tools are already in use across departments

- Often informally, sometimes with personal accounts
- Frequently without understanding risks or policies

What's at stake:

- Decisions depend on accurate information
- Public trust depends on transparent communications
- Public Records Act may apply to AI content
- Data breaches with PII create legal harm
- AI outputs can embed biases that affect decisions

What We'll Cover

- What GenAI is and how it differs from other tools
- Garden Grove's policy requirements
- Risks specific to GenAI
- When and how to use AI effectively
- Your responsibilities

Training manual: Contains detailed examples and scenarios

2. What is Generative AI

AI You Already Know

Most AI acts on existing data:

- Face recognition, medical imaging
- Speech recognition, spam detection
- Weather prediction, fraud detection, market analysis
- Self-driving cars

Generative AI creates new content:

- Text, images, audio, video
- Adapts to prompts, learns from conversation
- Common tools: ChatGPT, Claude, Gemini

Critical Distinction

GenAI is designed for creative and novel content generation, not necessarily factual or accurate content. This makes it powerful for brainstorming and drafting, but can be risky for decision-making and factual reporting if not grounded or double checked.

Understanding Terms (1 of 2)

Prompt: Your input to the AI

- Text you type (questions, instructions, context)
- Files you upload (documents, spreadsheets, images)
- Features you enable (web search, code execution)

Output: AI's response to your prompt

- Generated text, analysis, or recommendations
- Created visualizations or calculations
- Drafted documents or revised content

Understanding Terms (2 of 2)

Context: Information AI remembers within a conversation

- What you've discussed earlier in the same chat
- Gets "forgotten" when you start a new conversation

Training data: The massive collection of text AI learned from

- Books, websites, documents AI was trained on
- Determines what patterns AI recognizes

The basic flow: You provide a **prompt** → AI draws on its **training data** and conversation **context** → AI generates an **output** → You verify and refine the output for your needs.

3. Understanding the Risks

Risks not specific to GenAI, but can be easily overlooked

Hallucinations / Inaccurate Information

What it is: AI confidently generates false information, fabricated citations, invented statistics

Examples:

- Non-existent legal precedents in staff reports
- Fabricated budget numbers that "look right"
- Invented building code sections
- Fake demographic data for grants

Key point: AI is designed for creative content, not necessarily factual content

Data Loss / Privacy Concerns

What's at risk:

- Personnel actions memos
- Personally Identifiable Information
- Confidential negotiations
- Law enforcement data
- Infrastructure vulnerabilities
- Competitive bid details

How leaks happen:

- Prompts become training data
- Accidental sharing via history
- Personal/free accounts
- Malicious software in AI tools
- Malicious search sources

Bias / Discrimination

AI has no judgment, compassion, or reflection. They follow existing biases patterns in training data and potentially amplifies them.

- Job applicant screening may discriminate
- Budget recommendations may underserve neighborhoods
- Enforcement prioritization may reflect existing disparities
- Vendor selection may favor well-known companies

Why it happens:

- Training data overrepresents certain perspectives
- AI systems designed to be agreeable
- No built-in fairness or equity mechanism

Overreliance / Skill Atrophy

The risk:

- Staff stops developing evaluation expertise
- "The AI said so" becomes accepted justification
- Loss of institutional knowledge
- Critical thinking skills decline

Your expertise remains essential. Complacency and failure to critically evaluate AI output can be result of overreliance or skill atrophy.

Misuse / Over-confidence

Users may deliberately misuse AI to circumvent established procedures, or place excessive trust in AI recommendations for critical decisions.

Examples:

- AI-generated firewall config creates vulnerabilities
- Submitting AI staff reports without fact-checking
- Accepting AI legal interpretation without attorney review
- Using AI to find ways around procurement rules

4. Garden Grove's Gen AI Policy

Policy Framework

Foundation: AR 2.11 (Computer Systems) + AR 2.16 (Cloud Computing)

Core principle: Generative AI is a cloud computing service with additional protections due to specific risks

Key message: These aren't new restrictions - they extend existing policies you already follow

1. IT Authorization Required

Pre-approved tools:

- ChatGPT Enterprise
- Claude Enterprise

Why these:

- Better privacy/safety track record
- Enterprise versions don't train on your data

- Clear contractual protections

Why NOT others:

- Free versions may monetize your data
- Tools like Perplexity, DeepSeek, Meta AI, and X Grok have documented security issues

2. Account Requirements

Must:

- Use city-provisioned accounts only (IT creates)
- Use city email address
- Use multi-factor authentication
- Never share accounts

WARNING - Free Tier Accounts:

- Do NOT use free versions for city business
- Don't have free and enterprise accounts signed in simultaneously
- Easy to use wrong account by mistake

3. Data Protection

Never upload without IT Director + Department Director approval:

- Social security numbers
- Credit card numbers
- Payroll or personnel information
- Medical information
- Resident addresses or contact info
- Law enforcement data
- Contract details before public

All data created with city accounts = city property

4. Decision-Making and Accountability

AI is:

- A source of assistance
- A tool for drafting
- Pattern recognition

AI is NOT:

- A decision-maker
- Responsible for outputs
- Accountable for consequences

You are responsible for:

- All decisions made with AI assistance
- All outputs derived from AI
- Reviewing, revising, fact-checking all content

5. Bias Awareness

Use AI with understanding that it is **not objective or free of bias**

Apply extra scrutiny to things like:

- Hiring recommendations
- Service delivery decisions
- Enforcement prioritization
- Any decisions affecting people

6. Disclosure Requirements

Disclose AI use when transparency serves public interest:

When to disclose:

- AI analysis forms primary basis for recommendations
- AI tools interact with public (chatbots, translations)
- AI-generated images/video/audio could be mistaken for authentic
- AI analysis involved confidential or PII data

How to disclose:

- Clear and concise language
- Prominently placed
- Specific about AI's role

7. Public Records and Retention

Key points:

- Assume AI conversations = subject to Public Records Act
- Treat AI drafts like email drafts and working documents
- No formal retention schedule yet (in development)
- When in doubt: consult City Clerk

5. Understanding the Technology and the Human-AI Partnership

The Basics

AI is sophisticated pattern matching:

1. Trained on massive amounts of data (text, images, videos)
2. Learned statistical patterns in how words follow each other
3. Predicts most likely next word, then next, then next
4. Does this thousands of times per second

Analogy: If you've read thousands of staff reports, you recognize patterns and could draft one on an unfamiliar topic. AI does this but has "read" millions of documents and can spit out words really fast.

Why It Seems Capable

Patterns AI learned:

- How arguments are constructed
- How problems are typically solved
- Cause-and-effect relationships
- Different writing styles
- How to structure content types

This enables:

- Drafting coherent documents
- Summarizing texts
- Explaining concepts
- Analyzing data
- Generating variations

What AI Fundamentally Cannot Do

1. **No real understanding** - Example: recognizes patterns about building codes, doesn't understand why they exist or what makes structure safe
2. **No judgment or values** - can't determine if policy is ethical or fair
3. **No accountability** - no consequences when wrong
4. **No connection to current reality** - knowledge cutoff date, doesn't know yesterday's Council meeting
5. **No expertise or experience** - Example: has seen inspections described, hasn't conducted them. Can provide best practices, but hasn't learned from failures. No "feel" for when something's "off".

You remain essential

- You understand the actual context and constraints
- You provide judgment, ethics, and accountability
- You connect its outputs to real circumstances
- You bring expertise AI can only approximate

Understanding how GenAI works will allow you know how much to trust it, know its limitations, use it effectively as a tool and not treat it as an oracle.

The Human-AI Partnership

Humans provide:

- Critical thinking
- Ethics and equity
- Subject matter expertise
- Creativity and vision
- Accountability
- Context understanding
- Connection to current reality
- Common sense

AI provides:

- Fast information processing
- Pattern recognition
- Draft generation
- Format conversion
- Exploring variations
- Repetitive tasks
- Answers and how-tos

Questions to ask before using AI for a task

1. Do I understand the task well enough to evaluate output?

- If no: Learn manually first, then consider AI to augment you

2. Does this involve sensitive data?

- If yes: Do I have required approvals?

3. Does this require judgment, ethics, or accountability that AI can't provide?

- Policy decisions, equity considerations, anything affecting rights/services

6. Everyday Casual Queries

Most AI Use is Casual

Reality: Most people use AI like a search engine

- Quick questions
- Fast answers
- Move on

Totally fine! But recognize when to apply more rigor.

How AI Differs from Search Engines

Search Engines (Google, Bing)

- Shows multiple sources
- You see who's saying what
- You evaluate credibility
- See competing perspectives
- Links to original sources

AI Tools (ChatGPT, Claude)

- Synthesizes one answer
- Sources hidden even if you ask
- AI pre-selected what's credible
- Disagreements smoothed over

Why this matters

With search, you're doing the evaluation work yourself. With AI, the evaluation already happened invisibly. Sometimes that's fine. Sometimes it's risky.

Casual Usage Examples

Appropriate for:

- Definitions and explanations
- Format conversions
- Basic how-to questions, something you can verify right away
- Grammar and translation checks
- Brainstorming starting points

Refer to training manual for specific examples

When Casual Becomes Involved

Warning signs your "quick questions" needs more attention

- Answer will go in a document others will read
- You're making a decision based on the answer
- Involves resident/employee data or rights
- Policy interpretation or legal requirements
- Money or safety implications

When you feel you're in this territory: Shift to rigorous approach

Also Consider Traditional Methods

For code/policy/procedure questions:

- Check municipal code
- Check admin regulations (intranet)
- Ask supervisor or subject matter expert

For legal/regulatory questions:

- Consult City Attorney
- Check official state/federal agency sites
- Review actual statutes

7. Effective Prompting for Involved Tasks

Anatomy of an Effective Prompt

When casual becomes consequential, include:

1. **Task explanation:** What needs to be done
2. **Context:** Background, constraints, purpose
3. **Audience:** Who will read/use this
4. **Format:** How output should be structured
5. **Examples:** What good output looks like

See training manual for detailed prompt examples

Advanced Techniques (1 of 2)

Extended thinking/reasoning:

- Ask AI to show reasoning process
- Helps catch AI heading in wrong direction

Breaking large tasks:

- Execute complex projects in phases
- Start fresh conversations with focused prompts

Multiple versions:

- Request variations emphasizing different priorities
- Compare different approaches

Advanced Techniques (2 of 2)

Defining roles and perspective

- Use roles to shape focus and tone, not to add expertise

Know when to start fresh

- If a conversation has gone off track, starting a new conversation with a better prompt often works better than multiple corrections within the same conversation.

7B. Making AI Push Back

Why this matters

The problem: AI is designed to be agreeable

For consequential decisions, you need AI to:

- Question flawed premises
- Surface disagreement and complexity
- Acknowledge uncertainty
- Challenge your thinking

Solution: Custom instructions to combat agreeability

- See training manual for full instructions
- Helps surface bias and controversy

8. Evaluating AI Output

Why Evaluation Matters

The temptation:

- AI produces impressive-looking content quickly
- "It was right last time, why wouldn't it be right again?"
- Easy to accept as-is

The risk:

- AI can be confidently wrong

- Generates to feel helpful, not necessarily accurate

The requirement:

- You must verify facts
- You must ensure quality
- You must ensure ethical
- You must apply subject matter expertise

The Evaluation Process

1. Verify factual claims

- Check statistics against original sources
- Verify legal citations exist and say what AI claims
- Confirm dates, names, details

2. Check internal consistency

- Do paragraphs contradict each other?
- Do numbers add up?
- Are recommendations supported by analysis?

The Evaluation Process (continued)

3. Evaluate completeness

- Did AI address all aspects?
- Are there obvious gaps?

4. Assess bias and fairness

- Stereotypes or generalizations?
- Missing perspectives?

5. Question the scope

- Right scale of solution?
- Over-engineering something simple?

Set Aside Time

AI generates fast $\neq$ you should review fast

- Block time for verification
- Don't let AI speed pressure you
- This is like reviewing colleague's work
- Don't let AI replace your skill in finding/evaluating information

The Loop

When you find something wrong and not to your liking, keep working with AI to narrow it down:

- Specify what's wrong and why
- Provide examples of improvement
- Try different approaches

Then verify again

If you think AI is getting off track, ask it to write a summary prompt for you to start over in a new chat

9. Common Use Cases and Guidance

Use Case: Summarization

Good for: Meeting notes, long documents, email threads

Watch out for: May misidentify key points, miss context or subtext (limited attention)

Best practice: Use summary as preview, not substitute for understanding

Use Case: Question-Answer / Research

Good for: Finding info in documents you provide, explaining concepts

Watch out for: May invent citations, confidently answer questions it shouldn't

Best practice: Verify factual claims, disable web search if working with provided data only

Use Case: Drafting

Good for: First drafts, converting bullets to narrative, reformatting

Watch out for: Generic language, loss of your voice, skill atrophy

Best practice: Think of AI as intern who drafts but needs your expertise to finalize

Use Case: Data Analysis

Good for: Identifying patterns, creating visualizations

Watch out for: May misinterpret data, make calculation errors, find meaningless correlations

Best practice: Verify all calculations, provide context for interpreting patterns

Use Case: Brainstorming

Good for: Generating options, exploring approaches, challenging assumptions

Watch out for: Ideas may be impractical, lack understanding of constraints

Best practice: This is where AI excels - use freely, but apply judgment about feasibility

10. AI Tools: Web Search, Projects, and Code Execution

Built-in AI Tools

Modern enterprise AI platforms include:

1. Web Search
2. Projects
3. Code Execution / Data Analysis

These tools extend AI beyond text generation

Understanding them helps you get better results for specific tasks

Tool 1: Web Search

Web Search

What it does:

- Searches internet in real-time for current information
- Connects AI's knowledge to real, verifiable sources
- "Grounds" AI responses to reduce hallucinations

When to use it:

- Current events or recent news
- Recent policy/regulation changes
- Current data or statistics
- Information that changes frequently
- Verifying or grounding AI's claims

Web Search: Controls and Warnings

CRITICAL WARNING:

- Turn OFF web search when working with confidential data
- AI may use confidential data as search terms
- This leaks data to search engines and websites

How to control:

- Most enterprise tools let you toggle on/off
- Usually on by default
- May need to explicitly ask AI to activate it even if feature is enabled

Tool 2: Projects

What it does:

- Creates persistent virtual workspace
- AI accesses multiple documents across conversations
- Maintains custom instructions specific to project
- Remembers context accross converstaions

When to use it:

- Extended initiatives spanning multiple conversations
- Need AI to reference specific documents repeatedly
- Want consistent formatting across related work
- Want to define and reuse specific context
- Collaborative work with shared AI context

Project Setup Strategy

What to upload:

- Templates and examples
- Reference documents
- Style guides
- Standard language

Custom instructions:

- Define format preferences
- Set tone and style
- Specify constraints
- Establish workflow
- Describe context

Project Examples

See training manual for detailed setups:

1. Council Staff Reports Project

- Templates, examples, style guide uploaded
- Consistent formatting across all reports

2. Budget Analysis Project

- Current/prior year budgets, strategic plans
- All analysis references same data

3. Policy Development Project

- Research, comments, legal memos, drafts
- Maintain context through iterations

Project Best Practices

- Name projects clearly and specifically
- Only upload relevant documents
- Update documents as work evolves
- Set clear custom instructions
- Remove outdated documents
- Share projects across department when working together

Tool 3: Code Execution

Code Execution and Data Analysis

What it does:

- AI writes and runs code to perform tasks
- Analyzes data, creates visualizations
- Performs calculations, manipulates files
- Processes structured information
- Creates interactive tools and applications

Why use code for calculations:

- GenAI struggles with complex "mental math"
- Provides reliable calculations

Interactive Apps

Code execution creates interactive tools and applications:

- Web-based fee schedule calculators
- Interactive budget scenario planners
- Data visualization dashboards that update based on user inputs

AI builds the tool, you use it directly - no coding knowledge required

Code Execution: Use Cases

Data analysis:

- Spreadsheet/CSV analysis
- Charts and visualizations
- Statistical analysis
- Large dataset processing
- Format conversions

- Data cleaning

Interactive tools:

- Fee calculators
- Budget scenario planners
- Data dashboards
- Custom web-based tools

CRITICAL: Data Sensitivity Warning

When AI executes code on your data:

- Code may use external libraries connecting to outside services
- Data could be processed through third-party tools
- Some operations require internet connectivity
- Errors might expose data in logs

Best practice:

- **Only use code execution with non-confidential data**
- Treat with extra care around sensitive information

Code Execution Best Practices

- Double-check calculations, especially financial data
- Provide clean, well-organized data files
- Clearly specify needed calculations or analysis
- Request visualizations appropriate for audience
- Save generated charts/tables/tools (not persistent)
- Ask AI to revise if first attempt isn't right
- Request plain English explanation of code
- Always verify results
- When a simple spreadsheet would work just as well, probably skip this tool

You don't need to know how to code to use this effectively

Code Execution Examples

See training manual for detailed examples:

Recreation program attendance analysis:

- Calculate averages, participation rates
- Identify trends
- Generate comparison tables and charts

Interactive fee calculator:

- Takes program and age as input
- Applies discounts automatically
- Shows fee breakdown in real-time

Combining Tools

Example: Grant research and tracking

1. **Web search:** Find currently available grants
2. **Projects:** Create grant tracking project with requirements and deadlines
3. **Code execution:** Analyze budget data to match grant criteria

Strategic combination produces more powerful results

Key Takeaway: AI Tools

Web Search:

- Connects to current information
- Grounds responses in verifiable sources
- Turn OFF for confidential work

Code Execution:

- Reliable quantitative analysis
- Data processing and visualization
- Interactive tools
- Only use with non-confidential data

Projects:

- Persistent memory and context
- Consistent approach across work
- Shared knowledge base

Always with human oversight and verification

11. Taking Responsibility for AI Usage

Before Using AI

Ask yourself:

What data am I providing?

- PII, confidential, sensitive?
- Do I have approval?
- Am I sharing more than necessary?

Who has access?

- Using city-provisioned account?
- Free-tier service that trains on my data?
- What if there's a breach?

After Using AI

Ask yourself:

How will output be interpreted?

- Who will read this?
- Will they know AI was involved?
- Should disclosure be included?

What are ramifications of bad output?

- What decisions depend on this being accurate?
- Who is affected if wrong?
- Cost of error?

Professional Values Check

Does this align with:

- Your professional values?
- City's mission and goals?
- City policy?

Would you be comfortable:

- Explaining this process to Council?
- Explaining to residents?

Does this use respect the people affected?

12. Safety and Security Practices

Safety Practices

Checking citations:

- Be cautious with URLs - some may be malicious
- Don't provide personal data when checking citations
- Cross-check against authoritative sources

.

Disable web search for sensitive work:

- Turn OFF when working with confidential data you've provided
- Reduces risk of data exposure

AI Agents:

- Like screen sharing with AI system
- Take extra precautions
- Avoid sharing credentials

AI Tools to Avoid (Professionally and personally)

Avoid entirely:

- **Perplexity:** Known security vulnerabilities, monetizes your data
- **DeepSeek, Qwen, Doubao, Ernie:** Data jurisdiction concerns (China)
- **Meta AI, X Grok:** Designed to monetize your data like parent platforms
- **Any free-tier generative AI for work**

Why: No contractual protections, may train on your data, unclear data handling

Summary: Key Takeaways

Key Takeaways (1 of 2)

1. **AI is a tool, not a decision-maker** - You remain accountable
2. **Use only city-provisioned accounts** - Never personal/free AI accounts
3. **Protect sensitive data** - Don't upload PII or confidential info without approvals
4. **Be the subject matter expert** - You must evaluate AI outputs critically
5. **Verify everything** - Check facts, citations, calculations, logic

Key Takeaways (2 of 2)

- 6. **Disclose when transparency serves public interest** - Especially for major recommendations or public-facing work
- 7. **Watch for bias** - AI amplifies patterns in data, including discriminatory patterns
- 8. **Assume it's a public record** - AI conversations may be subject to PRA
- 9. **Stay human-centered** - Use AI to enhance work, not replace judgment

Questions?

Resources:

- Training manual (detailed examples and scenarios)
- IT helpdesk for account access
- City intranet for approved tools list
- City Attorney for legal questions
- City Clerk for records questions

Thank you for your attention